



## About

[About OptumSoft](#)[Fact Sheet](#)[FAQ](#)[Thought Leadership](#)[Executive Team](#)

## Legal

[Terms Of Use](#)[Privacy Policy](#)[Virtual Patent Marking](#)

## Products & Services

## Blog

[Jobs](#) Jan 5

## Contact

# Dangers of using dlsym() with RTLD\_NEXT

Posted In:

## Background

There are times when you want to wrap a library function in order to provide some additional functionality. A common example of this is wrapping the standard library's `malloc()` and `free()` so that you can easily track memory allocations in your program. While there are several techniques for wrapping library functions, one well-known [method](#) is using `dlsym()` with `RTLD_NEXT` to locate the *wrapped* function's address so that you can correctly forward calls to it.

## Problem

So what can go wrong? Let's look at an example:

## LibWrap.h

```
void* memAlloc(size_t s);  
// Allocate a memory block of size 's' bytes.  
void memDel(void* p);  
// Free the block of memory pointed to by 'p'.
```

## LibWrap.c

```
#define _GNU_SOURCE  
#include <dlfcn.h>  
#include "LibWrap.h"  
  
static void* malloc(size_t s) {  
    // Wrapper for standard library's 'malloc'.  
    // The 'static' keyword forces all calls to malloc() in this file to res  
    // to this functions.  
    void* (*origMalloc)(size_t) = dlsym(RTLD_NEXT,"malloc");  
    return origMalloc(s);  
}  
  
static void free(void* p) {  
    // Wrapper for standard library's 'free'.  
    // The 'static' keyword forces all calls to free() in this file to resol  
    // to this functions.  
    void (*origFree)(void*) = dlsym(RTLD_NEXT,"free");  
    origFree(p);  
}  
  
void* memAlloc(size_t s) {  
    return malloc(s);  
    // Call the malloc() wrapper.  
}  
  
void memDel(void* p) {  
    free(p);  
    // Call the free() wrapper.  
}
```

## Main.c

```
#include <malloc.h>  
#include "LibWrap.h"  
  
int main() {  
    struct mallinfo beforeMalloc = mallinfo();  
    printf("Bytes allocated before malloc: %d\n",beforeMalloc.uordblks);
```

```
void* p = memAlloc(57);
struct mallinfo afterMalloc = mallinfo();
printf("Bytes allocated after malloc: %d\n",afterMalloc.uordblks);

memDel(p);
struct mallinfo afterFree = mallinfo();
printf("Bytes allocated after free: %d\n",afterFree.uordblks);

return 0;
}
```

First compile LibWrap.c into a shared library:

```
$ gcc -Wall -Werror -fPIC -shared -o libWrap.so LibWrap.c
```

Next compile Main.c and link it against the libWrap.so that we just created:

```
$ gcc -Wall -Werror -o Main Main.c ./libWrap.so -ldl
```

Time to run the program!

```
$ ./Main
Bytes allocated before malloc: 0
Bytes allocated after malloc: 80
Bytes allocated after free: 0
```

So far, so good. No surprises. We allocated a bunch of memory and then freed it. The statistics returned by mallinfo() confirm this.

Out of curiosity, let's look at ldd output for the application binary we created.

```
$ ldd Main
linux-vdso.so.1 => (0x00007ffff1b1fe000)
./libWrap.so (0x00007fe7d2755000)
libdl.so.2 => /lib/x86_64-linux-gnu/libdl.so.2 (0x00007fe7d2542000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007fe7d217c000)
/lib64/ld-linux-x86-64.so.2 (0x00007fe7d2959000)
```



The ldd output is from Ubuntu 14.04.1 LTS for x86-64. Your output may differ.

Take note of the relative placement of libWrap.so with respect to libc.so.6: libWrap.so comes *before* libc.so.6. Remember this. It will be important later.

Now for fun, let's re-compile Main.c with libc.so.6 explicitly specified on the command-line and coming *before* libWrap.so:

```
$ gcc -Wall -Werror -o Main Main.c /lib/x86_64-linux-gnu/libc.so.6 ./libWrap
```

Re-run:

```
$ ./Main
Bytes allocated before malloc: 0
Bytes allocated after malloc: 80
Bytes allocated after free: 80
```

Uh oh, why are we leaking memory all of a sudden? We de-allocate everything we allocate, so why the memory leak?

It turns out that the leak is occurring because we are *not actually forwarding malloc() and free() calls to libc.so.6's implementations. Instead, we are forwarding them to malloc() and free() inside ld-linux-x86-64.so.2!*

“What are you talking about?!” you might be asking.

Well, it just so happens that ld-linux-x86-64.so.2, which is the dynamic linker/loader, has its own copy of malloc() and free(). Why? Because ld-linux has to allocate memory from the heap *before* it loads libc.so.6. But the version of malloc/free that ld-linux has does not actually free memory!



See elf/dl-minimal.c in glibc source code for ld-linux's malloc/free implementation.

But why does libWrap.so forward calls to ld-linux instead of libc? The answer comes down to how dlsym() searches for symbols when RTLD\_NEXT is specified. Here's the relevant excerpt from the dlsym(3) [man page](#):

| [RTLD\_NEXT] will find the next occurrence of a function in the search order after the current library. This allows one to provide a wrapper around a function in another shared library.


— dlsym(3)

To understand this better, take a look at ldd output for the new Main binary:

```
$ ldd Main
linux-vdso.so.1 => (0x00007fffe1da0000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f32c2e91000)
./libWrap.so (0x00007f32c2c8f000)
libdl.so.2 => /lib/x86_64-linux-gnu/libdl.so.2 (0x00007f32c2a8a000)
/lib64/ld-linux-x86-64.so.2 (0x00007f32c3267000)
```

Unlike [earlier](#), libWrap.so comes *after* libc.so.6. So when dlsym() is called inside libWrap.so to search for functions, it skips libc.so.6 since it precedes libWrap.so in the search order list. That means the searches continue through to ld-linux-x86-64.so.2 where

they find linker/loader's malloc/free and return pointers to those functions. And so, libWrap.so ends up forwarding calls to ld-linux instead of libc!

 Exercise to the reader: Verify that malloc/free calls are getting forwarded to ld-linux instead of libc by stepping through Main with GDB.

At this point you might be wondering: We ran a somewhat funky [command](#) to build our application and then encountered a memory leak due to weird library linking order caused by said command. Isn't this whole thing a silly contrived scenario?

The answer is unfortunately no. At OptumSoft, we recently encountered this very same memory leak with a binary compiled using the standard ./configure && make on x86-64 Ubuntu 14.04.1 LTS. For reasons we don't understand, the linking order for the binary was such that using dlsym() with RTLD\_NEXT to lookup malloc/free resulted in pointers to implementations inside ld-linux. It took a ton of effort and invaluable help from [Mozilla's rr tool](#) to root-cause the issue. After the whole ordeal, we decided to write a blog post about this strange behavior in case someone else encounters it in the future.

## Solution

If you find dlsym() with RTLD\_NEXT returning pointers to malloc/free inside ld-linux, what can you do?

For starters, you need to detect that a function address indeed does belong to ld-linux using dladdr():

```
void* func = dlsym(RTLD_NEXT, "malloc");
Dl_info dlInfo;
if(!dladdr(func, &dlInfo)) {
    // dladdr() failed.
}
if(strstr(dlInfo.dli_fname, "ld-linux")) {
    // 'malloc' is inside linker/loader.
}
```

Once you have figured out that a function is inside ld-linux, you need to decide what to do next. Unfortunately, there is no straightforward way to continue searching for the same function name in all other libraries. But if you know the name of a specific library in which the function exists (e.g. libc), you can use dlopen() and dlsym() to fetch the desired pointer:

```
void* handle = dlopen("libc.so.6", RTLD_LAZY);
// NOTE: libc.so.6 may *not* exist on Alpha and IA-64 architectures.
if(!handle) {
```

```
// dlopen() failed.
}
void* func = dlsym(handle, "free");
if(!func) {
    // Bad! 'free' was not found inside libc.
}
```



dlopen'ing a library to replace malloc/free is generally frowned upon. Use at your own risk.

## Summary

- One can use `dlsym()` with `RTLD_NEXT` to implement wrappers around `malloc()` and `free()`.
- Due to unexpected linking behavior, `dlsym()` when using `RTLD_NEXT` can return pointers to `malloc/free` implementations inside `ld-linux` (dynamic linker/loader). Using `ld-linux`'s `malloc/free` for general heap allocations leads to memory leaks because that particular version of `free()` doesn't actually release memory.
- You can check if an address returned by `dlsym()` belongs to `ld-linux` via `dladdr()`. You can also lookup a function in a specific library using `dlopen()` and `dlsym()`.

« [The TACC® Difference](#)

### Post A Comment

Write your message here...

Your Full Name

E-mail Address

Website

Submit

Search

Start Typing

© OptumSoft 2017

© OptumSoft 2017